

# Deep learning

Alexandre Boulch

[www.boulch.eu](http://www.boulch.eu)

IOGS - ATSI


# Outline

- Back on last course
- Concept of deep learning
- Convolutional Neural Networks
- Attention and transformers

# Back on neural networks

## The linear layer


- Also called *fully connected*
  - a neuron is connected to all the inputs
- High number of parameters (**W** matrix):  
 $|inputs| * |outputs|$




# Multi-layer perceptron

A stack of linear layers with activation functions (e.g., sigmoids)

Optimization with **gradient descent**.



# Optimization: forward-backward algorithm



# Chain rule applied to neural networks

**Forward**



# Chain rule applied to neural networks

**Forward**


$$y_1 = f_1(x_n, \theta_1) \quad \rightarrow \quad y_2 = f_2(y_1) \quad \rightarrow \quad y_3 = f_3(y_2, \theta_2) \quad \rightarrow \quad \hat{y} = L(y_3, y_n)$$

**Backward**

$$\frac{\partial L}{\partial \theta_1} = \frac{\partial f_1}{\partial \theta_1} \frac{\partial L}{\partial y_1} \leftarrow \frac{\partial L}{\partial y_1} = \frac{\partial f_2}{\partial y_1} \frac{\partial L}{\partial y_2} \leftarrow \frac{\partial L}{\partial y_2} = \frac{\partial f_3}{\partial y_2} \frac{\partial L}{\partial y_3} \leftarrow \frac{\partial L}{\partial y_3} = \dots$$

$$\frac{\partial L}{\partial \theta_2} = \frac{\partial f_3}{\partial \theta_2} \frac{\partial L}{\partial y_3}$$


# Chain rule applied to neural networks



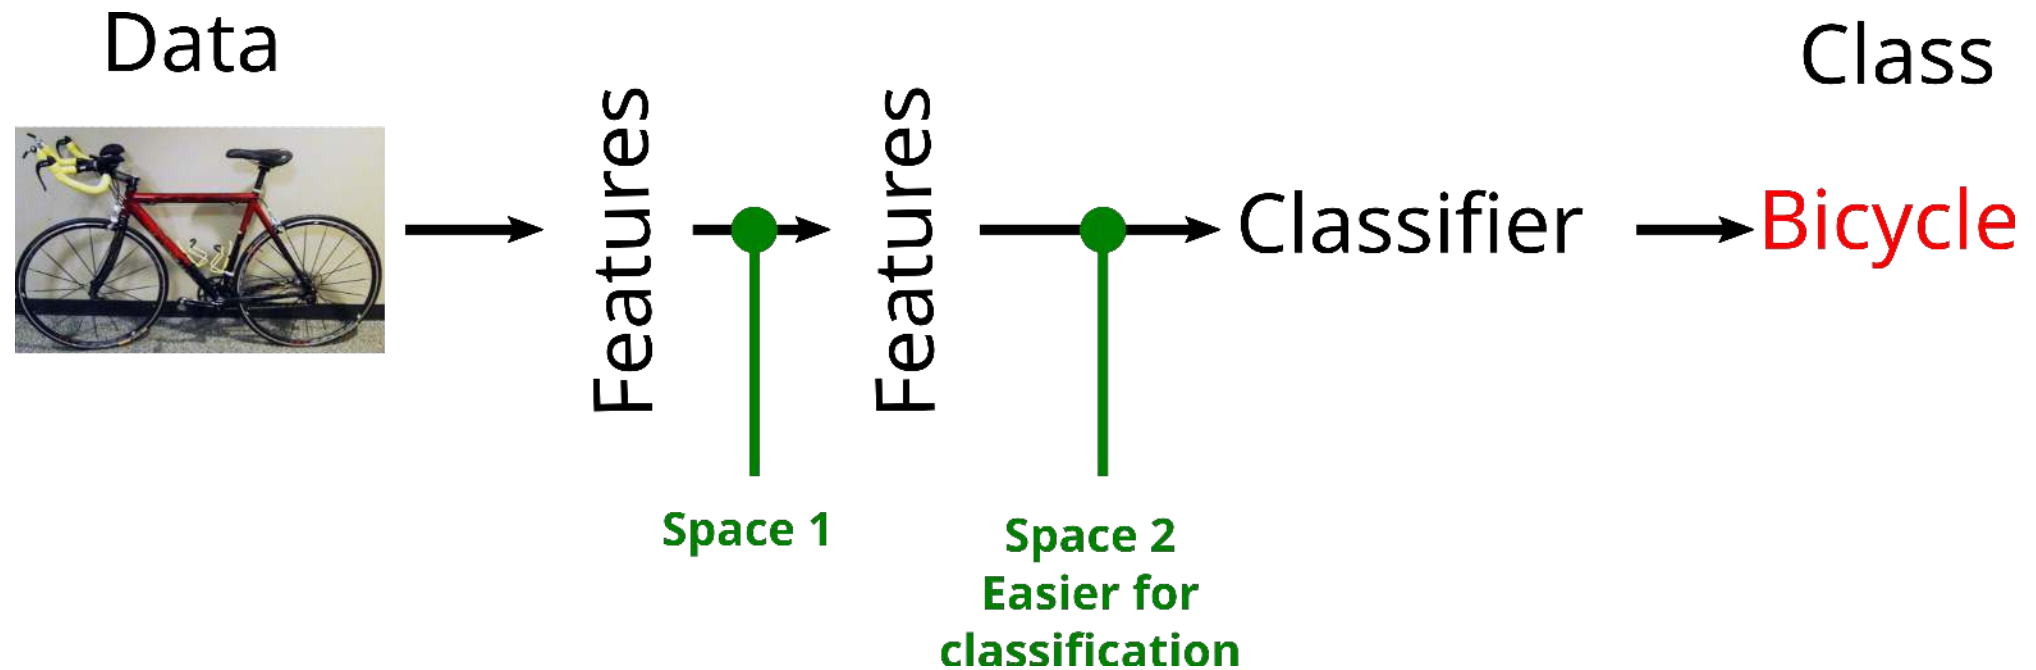


# Deep learning concept


# Deep learning concept




# Deep learning concept



# Deep learning concept



# Massively data driven approaches



# Convolutions and image processing


# Multi-layer perceptron (before 1990)

## MLP becomes larger and deeper

- difficult convergence
- few data
- very long training
- progressive loss of interest

## SVM


- simple to use
- convergence proof
- fast



# Multi-layer perceptron for images

## Using a linear layer ?

Lots of weights! (at least one per pixel!)






# Multi-layer perceptron for images

## Using a linear layer ?

Lots of weights! (at least one per pixel!)


Is it interesting to look at relations in the whole image ?



Look at  
the whole  
image

# Convolution

Look at small neighborhoods (where the objects are)




Look at  
neighborhoods

# Convolution

Look at small neighborhoods (where the objects are)


Create neurons that take a patch input



# Convolution

## Problem

Translation of the object must lead to same behaviour of the neurons



Same behaviour for the same pattern at different location


# Convolution

## Problem

Translation of the object must lead to same behaviour of the neurons

## Solution

Use the same neuron (i.e. all the neurons of the layer share weights)



# Convolution


## Forward

Let  $(i, j)$  be the coordinates in the input map.

$(k, l)$  be the size of the patch (size of the kernel, usually  $k = l$ )

Then:

$$y_{i,j} = \sum_k \sum_l w_{k,l} x_{i+k,j+l} + b$$



# Convolution

Backward weight update

$$\frac{\partial y_{i,j}}{\partial w_{k,l}} = x_{i+k,j+l}$$

Let  $y$  be the output map and  $\Delta y$  be the gradient coming back:

$$\frac{\partial y}{\partial w_{k,l}} = \sum_i \sum_j x_{i+k,j+l}$$

Finally, the update rule:


$$\begin{aligned} w_{k,l} &\leftarrow w_{k,l} - \alpha \frac{\partial y}{\partial w_{k,l}} \Delta y \\ &\leftarrow \sum_i \sum_j x_{i+k,j+l} \Delta y_{i,j} \end{aligned}$$

# Convolutions

## Forward


Same with term to term multiplication:

$$y_{i,j} = \sum_a \sum_b \mathbf{w}_{a,b} \mathbf{x}_{i+a,j+b}$$





# Convolution: what do convolutions learn?



Gabor filters.



First layer of AlexNet.

# Dimension reduction


With the previous convolution, the output dimension is the same as the input dimension.

For classification: only one label, need for `\textbf{dimension reduction}`.

- **convolution stride**: do not look at all the pixels of the input (one every two, one every three...)
- **Max Pooling**

# Max Pooling

- dimension reduction
- relative translation invariability




# Max Pooling

## Forward

Max signal


## Backward

Gradient transmission to max  
signal origin, zero otherwise



# Convolutional Neural Networks - LeNet (1990)

LeNet (1990)  
Images 28x28



Very good results on digits recognition !

# Issues

## Issues

- Learning speed
- Exploding or vanishing gradients
- Overfitting
- Local minima


## Limitations

- Architecture
- Initialization
- Computing power
- Data
- Optimization


# Solutions

- activations
- mini-batches
- batch norm
- good weight initialization
- better optimization


# Activations



Tangente  
hyperbolique



Sigmoïde





Rectified Linear  
Unit (ReLU)

## Rectified linear unit}

- Faster gradient computation
- Similar convergence




# Mini-batches



## Gradient smoothing

Smoother gradient converges faster.

# BatchNorm



Changes in the signal dynamic make the model more difficult to optimize: exponential or vanishing gradients.

**Objective:** control the signal distribution:

$$y^{l*} = \frac{y^l - \mu}{\sqrt{\sigma^2 + \epsilon}} \gamma + \beta$$

$\gamma$  and  $\beta$  are learnt,  $\mu$  and  $\sigma$  are computed (mean and standard deviation).

Learning is faster (iteration number) but slower (statistics computation).

# Weight initialization}

Weights have great influence on convergence speed.

They are randomly initialized.

- too small weights: vanishing signal
- too high: exploding signal

Conservation of signal properties.

$$\text{Var}(Y) = \text{Var}(X)$$

# Weight initialization

## Xavier initialization

$X \in \mathbb{R}^n$ , weights  $W$  and output  $Y \in \mathbb{R}$

$$Y = W_1X_1 + W_2X_2 + \dots + W_nX_n$$

$X_i$  and  $W_i$  independent:

$$\text{Var}(W_iX_i) = E[X_i]^2\text{Var}(W_i) + \text{Var}(X_i)\text{Var}(W_i) + \text{Var}(X_i)E[W_i]^2$$

$E[X_i] = 0$  and  $E[W_i] = 0$ :

$$\text{Var}(W_iX_i) = \text{Var}(X_i)\text{Var}(W_i)$$


finally \$ \text{Var}(Y) = n \text{Var}(X\_i) \text{Var}(W\_i)\$ and we chose

$$\text{Var}(W_i) = \frac{1}{n}$$

# Optimization - Stochastic Gradient Descent

See neural network class

Update rule:  $w_{t+1} = w_t + \alpha \Delta w$  (learning rate  $\alpha$ )



- Step decrease
- Exponential decrease
- Cosine annealing ...

# Optimization - SGD with Momentum

Same idea as mini batch: smooth gradient in the good direction


## Momentum

Use previous gradient to ponderate the direction of the new gradient.


$$v_t = \gamma v_{t-1} + \alpha \Delta w$$

$$w_t = w_{t-1} - v_t$$

$\gamma$  is the momentum.



# ResNet



**Do not forget the classics**




# Data

- Representative
- Data augmentation
- Data normalization

# Data

Train data must be representative of the problem}



# Data

## Data normalization


- Compute mean  $\mu$  and standard deviation  $\sigma$  on the train set.
- Normalize input  $I$  (train and test):

$$\hat{I} = \frac{I - \mu}{\sigma}$$

# Data - data augmentation

Random variations of input parameters (images: lightness, contrast \dots)

- train on a more representative set
- avoid learning on unwanted features



# Problems and partial solutions

## Problems

- Small amount of data
- Low computational power

## Solutions ?

- Use classical approaches (Perceptron, SVM, ...)